

Московский государственный технический университет им. Н.Э. Баумана
Факультет «Информатика и системы управления»
Кафедра «Системы обработки информации и управления»



**«Методы машинного обучения в
автоматизированных системах обработки
информации и управления»
Лабораторная работа №7
«Алгоритмы Actor-Critic»**

ИСПОЛНИТЕЛЬ:

Демирев Н.К.
Группа ИУ5-21М

" " _____ 2023 г.

Москва 2023

1. Задание

- Реализуйте любой алгоритм семейства Actor-Critic для произвольной среды.

2. Листинг

2.1. policy.py

```
import torch.nn as nn
import torch.nn.functional as F

class Policy(nn.Module):
    def __init__(self):
        super(Policy, self).__init__()
        self.affine1 = nn.Linear(6, 128)

        # actor's layer
        self.action_head = nn.Linear(128, 3)

        # critic's layer
        self.value_head = nn.Linear(128, 1)

        # action & reward buffer
        self.saved_actions = []
        self.rewards = []

    def forward(self, x):
        x = F.relu(self.affine1(x))

        # actor: choses action to take from state s_t
        # by returning probability of each action
        action_prob = F.softmax(self.action_head(x), dim=-1)

        # critic: evaluates being in the state s_t
        state_values = self.value_head(x)

        # return values for both actor and critic as a tuple of 2 values:
        # 1. a list with the probability of each action over the action space
        # 2. the value from state s_t
        return action_prob, state_values
```

2.2. main.py

```
import gymnasium as gym
import numpy as np
from itertools import count
from collections import namedtuple
```

```

import torch
import torch.nn.functional as F
import torch.optim as optim
from torch.distributions import Categorical
from policy import Policy
import os
import pygame
from tqdm import tqdm

os.environ['SDL_VIDEODRIVER'] = 'dummy'
pygame.display.set_mode((640, 480))

# Cart Pole
CONST_ENV_NAME = 'Acrobot-v1'
env = gym.make(CONST_ENV_NAME)
GAMMA = 0.99
SavedAction = namedtuple('SavedAction', ['log_prob', 'value'])

model = Policy()
optimizer = optim.AdamW(model.parameters(), lr=1e-3)
eps = np.finfo(np.float32).eps.item()

def select_action(state):
    state = torch.from_numpy(state).float()
    probs, state_value = model(state)

    # create a categorical distribution over the list of probabilities of actions
    m = Categorical(probs)

    # and sample an action using the distribution
    action = m.sample()

    # save to action buffer
    model.saved_actions.append(SavedAction(m.log_prob(action), state_value))

    # the action to take (left or right)
    return action.item()

def finish_episode():
    """
    Training code. Calculates actor and critic loss and performs backprop.
    """
    R = 0
    saved_actions = model.saved_actions
    policy_losses = [] # list to save actor (policy) loss
    value_losses = [] # list to save critic (value) loss
    returns = [] # list to save the true values

    # calculate the true value using rewards returned from the environment

```

```

for r in model.rewards[::-1]:
    # calculate the discounted value
    R = r + GAMMA * R
    returns.insert(0, R)

returns = torch.tensor(returns)
returns = (returns - returns.mean()) / (returns.std() + eps)

for (log_prob, value), R in zip(saved_actions, returns):
    advantage = R - value.item()

    # calculate actor (policy) loss
    policy_losses.append(-log_prob * advantage)

    # calculate critic (value) loss using L1 smooth loss
    value_losses.append(F.smooth_l1_loss(value, torch.tensor([R])))

# reset gradients
optimizer.zero_grad()

# sum up all the values of policy_losses and value_losses
loss = torch.stack(policy_losses).sum() + torch.stack(value_losses).sum()

# perform backprop
loss.backward()
optimizer.step()

# reset rewards and action buffer
del model.rewards[:]
del model.saved_actions[:]

def main():
    running_reward = -500

    # run infinitely many episodes
    for i_episode in count(1):
        # print(running_reward)
        # reset environment and episode reward
        state, _ = env.reset()
        ep_reward = 0
        # for each episode, only run 9999 steps so that we don't
        # infinite loop while learning
        for t in range(1, 9999):
            # select action from policy
            action = select_action(state)
            # take the action
            state, reward, done, truncated, _ = env.step(action)
            model.rewards.append(reward)
            ep_reward += reward
            if done or truncated:
                break

```

```

    # print(ep_reward)
    # update cumulative reward
    running_reward = 0.05 * ep_reward + (1 - 0.05) * running_reward
    # perform backprop
    finish_episode()
    # log results
    if i_episode % 10 == 0:
        print(f"Episode {i_episode}\tLast reward: {ep_reward:.2f}\tAverage
reward: {running_reward:.2f}")
    # check if we have "solved" the cart pole problem
    if running_reward > env.spec.reward_threshold * 2:
        print(f"Solved! Running reward is now {running_reward} and the last
episode runs to {t} time steps!")
        break

env2 = gym.make(CONST_ENV_NAME, render_mode='human')
# reset environment and episode reward
state, _ = env2.reset()
ep_reward = 0
# for each episode, only run 9999 steps so that we don't
# infinite loop while learning
bar = tqdm(range(1, 10000), bar_format=' {l_bar}{bar:20}{r_bar}{bar:-10b}',
colour='CYAN')
for t in bar:
    # select action from policy
    action = select_action(state)
    # take the action
    state, reward, done, _, _ = env2.step(action)
    model.rewards.append(reward)
    ep_reward += reward
    if done:
        bar.update(10000-t)
        bar.refresh()
        bar.close()
        break

if __name__ == '__main__':
    main()

```

3. Экранные формы

```
(venv) PS G:\repos\MMO\7 lab> python.exe .\main.py
Episode 10      Last reward: -500.00    Average reward: -496.82
Episode 20      Last reward: -500.00    Average reward: -490.92
Episode 30      Last reward: -327.00    Average reward: -474.30
Episode 40      Last reward: -445.00    Average reward: -466.38
Episode 50      Last reward: -357.00    Average reward: -423.47
Episode 60      Last reward: -487.00    Average reward: -385.33
Episode 70      Last reward: -240.00    Average reward: -347.07
Episode 80      Last reward: -214.00    Average reward: -294.96
Episode 90      Last reward: -203.00    Average reward: -247.62
Episode 100     Last reward: -154.00    Average reward: -221.24
Episode 110     Last reward: -173.00    Average reward: -209.70
Episode 120     Last reward: -149.00    Average reward: -202.95
Solved! Running reward is now -198.17254604049728 and the last episode runs to 157 time steps!
100%|████████████████████| 9999/9999 [00:12<00:00, 822.84it/s]
(venv) PS G:\repos\MMO\7 lab> 
```